

# ECGR 5105 Homework 6: Neural Networks

Lauren Bourque (Student ID: 801264851)

November 26, 2023

## GitHub Link

[Click here to view the code](#)

## Problem 1a.

### Training and Validation Loss Results:

The loss values for the regression from Homework 5 with a learning rate of 0.001 can be seen below.

Figure 1: Homework 5 Loss Values

Learning rate of 0.001		
Epochs: 500,	Training Loss: 0.634329,	Validation Loss: 0.638253
Epochs: 1000,	Training Loss: 0.340199,	Validation Loss: 0.411084
Epochs: 1500,	Training Loss: 0.319608,	Validation Loss: 0.394241
Epochs: 2000,	Training Loss: 0.316399,	Validation Loss: 0.388427
Epochs: 2500,	Training Loss: 0.315724,	Validation Loss: 0.385703
Epochs: 3000,	Training Loss: 0.315560,	Validation Loss: 0.384277
Epochs: 3500,	Training Loss: 0.315516,	Validation Loss: 0.383491
Epochs: 4000,	Training Loss: 0.315504,	Validation Loss: 0.383048
Epochs: 4500,	Training Loss: 0.315501,	Validation Loss: 0.382797
Epochs: 5000,	Training Loss: 0.315500,	Validation Loss: 0.382654

The loss values for the neural network can be seen below. As can be seen from the results from Figure 2 compared to the results from Figure 1, the neural network achieves an eventual better training loss but slightly worse validation loss than the linear regression model from Homework 5. The neural network has  $(12 \times 32 + 32) + (32 \times 1 + 1) = 417$  trainable parameters compared to the 12 parameters from the linear regression model.

Figure 2: Neural Network Loss Values

Epoch 0/5000,	Training Loss: 0.8740482330322266,	Validation Loss: 1.2669895887374878
Epoch 500/5000,	Training Loss: 0.270549476146698,	Validation Loss: 0.4624500572681427
Epoch 1000/5000,	Training Loss: 0.285550057888031,	Validation Loss: 0.48047706484794617
Epoch 1500/5000,	Training Loss: 0.30756011605262756,	Validation Loss: 0.47917163372039795
Epoch 2000/5000,	Training Loss: 0.39476126432418823,	Validation Loss: 0.47147655487060547
Epoch 2500/5000,	Training Loss: 0.24433428049087524,	Validation Loss: 0.4641323387622833
Epoch 3000/5000,	Training Loss: 0.09919007867574692,	Validation Loss: 0.4574287533760071
Epoch 3500/5000,	Training Loss: 0.30647051334381104,	Validation Loss: 0.4536994397640228
Epoch 4000/5000,	Training Loss: 0.15413716435432434,	Validation Loss: 0.44924670457839966
Epoch 4500/5000,	Training Loss: 0.37633734941482544,	Validation Loss: 0.44976645708084106
Epoch 5000/5000,	Training Loss: 0.25254371762275696,	Validation Loss: 0.44664818048477173

## Problem 1b.

### Training and Validation Loss Results:

The loss values for the neural network with 3 hidden layers can be seen below.

Figure 3: Neural Network with 3 Layers Loss Values

```
Epoch 0/5000, Training Loss: 1.186692237854004, Validation Loss: 1.1106796264648438
Epoch 500/5000, Training Loss: 0.26614847779273987, Validation Loss: 0.41631677746772766
Epoch 1000/5000, Training Loss: 0.571448028087616, Validation Loss: 0.46390071511268616
Epoch 1500/5000, Training Loss: 0.3825329542160034, Validation Loss: 0.468228816986084
Epoch 2000/5000, Training Loss: 0.07898759841918945, Validation Loss: 0.4665388762950897
Epoch 2500/5000, Training Loss: 0.10419765114784241, Validation Loss: 0.4629688858985901
Epoch 3000/5000, Training Loss: 0.08431097865104675, Validation Loss: 0.4802996814250946
Epoch 3500/5000, Training Loss: 0.19947503507137299, Validation Loss: 0.5038023591041565
Epoch 4000/5000, Training Loss: 0.04506227746605873, Validation Loss: 0.530514121055603
Epoch 4500/5000, Training Loss: 0.08443551510572433, Validation Loss: 0.5555590391159058
Epoch 5000/5000, Training Loss: 0.07330196350812912, Validation Loss: 0.5717246532440186
```

The neural network with 3 hidden layers has a much better training loss than both the neural network with 1 hidden layer and the linear regression model. However, it has a worse validation loss than both the neural network with 1 hidden layer and the linear regression model. There are  $(12 \times 32 + 32) + (32 \times 64 + 64) + (64 \times 16 + 16) + (16 \times 1 + 1) = 9409$  parameters versus the 417 trainable parameters from problem 1a and the 12 trainable parameters from the linear regression model.

## Problem 2a.

### Training Time, Training Loss, and Evaluation Accuracy Results:

The total training time was 8241 seconds or 2.289 hours.

The training loss was 0.5180834466234192 at epoch 291.

The final evaluation accuracy was 49.72%, but the best recorded accuracy was 50.49%

## Problem 2b.

### Training Time, Training Loss, and Evaluation Accuracy Results:

The total training time was 8654 seconds or 2.404 hours.

The training loss was 0.009872898049152378 at epoch 291.

The final evaluation accuracy was 51.40%, but the best recorded accuracy was 54.01%

The model had a better accuracy than the model from problem 2a, showing that adding more hidden layers improved the accuracy.

The model in 2a had  $(3072 \times 512 + 512) + (512 \times 10 + 10) = 1578506$  parameters, while the model in 2b had  $(3072 \times 512 + 512) + (512 \times 256 + 256) + (256 \times 128 + 128) + (128 \times 10 + 10) = 1748890$  parameters.